

CMPUT 567: Machine Learning 2

Keven Qiu
Instructor: Martha White

1

The generalization error for a given f , with squared error is

$$E_{x,y}[(f(X) - Y)^2] = \underbrace{E[(f(X) - f^*(X))^2]}_{\text{reducible error}} + \underbrace{E[(f^*(X) - Y)^2]}_{\text{irreducible error}}$$

where $f^*(x) = E[Y|x]$ and the irreducible error is $E_x[\text{Var}(Y|X)]$.

Bias and variance is about the function class \mathcal{F} (and algorithm, i.e. objective and optimizer).

$$f_{\mathcal{D}} = \text{Alg}(\mathcal{D})$$

where $\mathcal{D} \sim \underbrace{p(x_1, y_1)p(x_2, y_2) \cdots p(x_n, y_n)}_{p(\mathcal{D})}$. Then

$$E_{\mathcal{D}}[GE(f_{\mathcal{D}})] = \underbrace{E_{\mathcal{D}}[\text{reducible error } (f_{\mathcal{D}})]}_{\text{bias}^2 + \text{variance}} + \text{irreducible error}$$

Note there are two variances: the irreducible error and the variance of our function f .

For one x :

$$\begin{aligned} E_{\mathcal{D}}[(f_{\mathcal{D}}(x) - f^*(x))^2] &= E_{\mathcal{D}}[(\underbrace{f_{\mathcal{D}}(x) - E_{\mathcal{D}}[f_{\mathcal{D}}(x)]}_a + \underbrace{E_{\mathcal{D}}[f_{\mathcal{D}}(x)] - f^*(x)}_b)^2] \\ &= \underbrace{b^2}_{\text{bias}^2} + \underbrace{a^2}_{\text{variance}} + 2 \underbrace{E_{\mathcal{D}}[f_{\mathcal{D}}(x) - E_{\mathcal{D}}[f_{\mathcal{D}}(x)]]}_b b \\ &= (E_{\mathcal{D}}[f_{\mathcal{D}}(x)] - f^*(x))^2 + \text{Var}_{\mathcal{D}}(f_{\mathcal{D}}(x)) \end{aligned}$$

If we can represent f^* (high model complexity), then $E_{\mathcal{D}}[f_{\mathcal{D}}(x)] \approx f^*(x)$ for all x .

1.1 Bias and Variance for Linear Regression

With nonlinear representation Φ , we have by SVD

$$\Phi = U\Sigma V^T \in \mathbb{R}^{n \times p}$$

where $U = [u_1, \dots, u_n]$, with left singular vectors $u_j \in \mathbb{R}^{n \times 1}$ and $V = [v_1, \dots, v_p]$, with right singular vectors $v_j \in \mathbb{R}^{p \times 1}$, and $\Sigma \in \mathbb{R}^{n \times p}$.

We consider thin SVD, where we consider only the first p rows/columns.

$$\Sigma = \begin{bmatrix} \sigma_1 & \cdots & 0 \\ 0 & \ddots & 0 \\ 0 & \cdots & \sigma_p \\ \vdots & \vdots & \vdots \\ 0 & \cdots & 0 \end{bmatrix}, \Sigma_p = \text{diag}(\sigma_1, \dots, \sigma_p), U_p = [u_1, \dots, u_p]$$

since $U\Sigma = U_p\Sigma_p$. Note that $\Sigma^T\Sigma = \Sigma_p^2$.

$$\begin{aligned} \Phi^T\Phi &= V\Sigma^T U^T U\Sigma V^T \\ &= V\Sigma^T \Sigma V^T \\ &= V\Sigma_p^2 V^T \\ \mathbf{w} &= (\Phi^T\Phi)^{-1}\Phi^T y \\ &= V\Sigma_p^{-2} V^T V\Sigma_p U_p^T y \\ &= V\Sigma_p^{-1} U_p^T y \\ &= \sum_{j=1}^p \frac{u_j^T y}{\sigma_j} v_j \end{aligned}$$

Note that for this, all singular values must be nonzero/positive. $\Phi^t = V\Sigma_p^{-1}U_p^T$ is known as the pseudoinverse of Φ .

If Φ has lower rank $r < p$, i.e. $\sigma_1, \dots, \sigma_r > 0, \sigma_{r+1}, \dots, \sigma_p = 0$, then $\Phi^t = V_r\Sigma_r^{-1}U_r^T$ where $V_r = [v_1, \dots, v_r]$.

Notice that \mathbf{w} is sensitive to small singular values since we divide by them. Reasons why σ_j can be small:

- linearly correlated features (one dimension might not have anything)
- insufficient data

1.2 ℓ_2 -Regularization

For linear regression, the MLE is $p(y|x) = \mathcal{N}(\phi w, \sigma^2)$. MAP with Gaussian prior is $-\ln p(w_j) = \text{constant} + \frac{\lambda}{2\sigma^2} w_j^2$.

$$c(w) = \frac{1}{2} \|\Phi w - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2 \implies (\Phi^T\Phi + \lambda I)w = \Phi^T y \implies w = (\Phi^T\Phi + \lambda I)^{-1}\Phi^T y$$

Using SVD $\Phi^T\Phi + \lambda I = V(\Sigma_p^2 + \lambda I)V^T$,

$$w_{MAP} = V(\Sigma_p^2 + \lambda I)^{-1}U_p^T y = \sum_{j=1}^p \frac{\sigma_j}{\sigma_j^2 + \lambda} (u_j^T y) v_j$$

As long as large λ , the fraction does not blow up like it did without ℓ_2 -regularization.

1.3 Better Generalization Error

Realizable setting: $y \sim \mathcal{N}(\Phi w^*, \sigma^2)$ for some w^* and $f^* \in \mathcal{F}$ (f^* is in function class).

Let $w(\mathcal{D})$ be our estimator.

$$MSE(w(\mathcal{D})) = E_{\mathcal{D}} \left[\sum_{j=1}^p (w_j(\mathcal{D}) - w_j^*)^2 \right] = \sum_{j=1}^p (E[w_j(\mathcal{D})] - w_j)^2 + \sum_{j=1}^p \text{Var}(w_j(\mathcal{D}))$$

Notice $y = \phi w^* + \varepsilon$ where $\varepsilon \sim \mathcal{N}(0, \sigma^2)$. For MLE,

$$\begin{aligned} E[w_{MLE}(\mathcal{D})] &= E[(\Phi^T \Phi)^{-1} \Phi^T y] \\ &= E[(\Phi^T \Phi)^{-1} \Phi^T (\Phi w^* + \varepsilon)] \\ &= E[(\Phi^T \Phi)^{-1} (\Phi^T \Phi) w^*] + \underbrace{E[(\Phi^T \Phi)^{-1} \Phi^T]}_{=0} E[\varepsilon] \\ &= w^* \end{aligned}$$

This is unbiased as it equals w^* . For MLE variance,

$$\sum_{j=1}^p \text{Var}(w_{MLE}(\mathcal{D})) = \sigma^2 E \left[\sum_{j=1}^p \sigma_j^{-2} \right]$$

where first σ^2 is variance, not singular value. This can be very big if σ_p is often small.

For MAP,

$$E[w_{MAP}(\mathcal{D})] = E[(\Phi^T \Phi + \lambda I)^{-1} \Phi^T (\Phi w^* + \varepsilon)] \neq w^*$$

so this is biased as λ increases.

$$\sum_{j=1}^p \text{Var}(w_{MAP}(\mathcal{D})) = \sigma^2 E \left[\sum_{j=1}^p \frac{\sigma_j^2}{(\sigma_j^2 + \lambda)^2} \right]$$

This is not very big if λ is large.

Non-realizable: $y = \mathcal{N}(f^*(x), \sigma^2)$, but $f^* \notin \mathcal{F}$. w_{MLE} is also biased.

Note about n : We can stick a $\frac{1}{n}$ in front of $\Phi^T \Phi$ to act like an average over outerproducts. Note $\frac{1}{n} \Phi^T \Phi = E[\Phi^T \Phi]$.

Without normalization, $\sigma_j^2 \rightarrow \infty$ as $n \rightarrow \infty$. With normalization $\frac{1}{n} \sigma_j^2 \rightarrow$ singular values for $E[\Phi^T \Phi]$.

$\sum_{j=1}^p \sigma_j^{-2} \rightarrow 0$ as $n \rightarrow \infty$, so $\text{Var}(w_{MLE}(\mathcal{D})) \rightarrow 0$ as $n \rightarrow \infty$.

2

Multiple outputs: $y \in \mathbb{R}^m, \phi(x)W \approx y$.

$$Y = \begin{bmatrix} y_{11} & \cdots & y_{1m} \\ \vdots & \ddots & \vdots \\ y_{n1} & \cdots & y_{nm} \end{bmatrix} = \begin{bmatrix} \mathbf{y}_1 & \cdots & \mathbf{y}_m \end{bmatrix}$$

where $\mathbf{y}_i \in \mathbb{R}^n$. So

$$\mathbf{w} = \Phi^t Y = \begin{bmatrix} \Phi^t \mathbf{y}_1 & \cdots & \Phi^t \mathbf{y}_m \end{bmatrix}$$

2.1 Optimization

Find

$$\min_{w \in \mathbb{R}^p} c(w)$$

For example, $c(w) = \frac{1}{n} \sum_{i=1}^n c_i(w)$, linear regression $c_i(w) = \frac{1}{2}(\phi(x_i)w - y_i)^2$ or logistic regression $c_i(w) = -y_i \ln f_w(x_i) - (1 - y_i) \ln(1 - f_w(x_i))$.

1. Find a stationary point w_0 , $\nabla c(w_0) = 0$.
2. Check type using second derivative test ($c'_i(w_t) < 0, > 0, = 0$).

In p -dimensions, the **Hessian** matrix $H_{c(w_t)} \in \mathbb{R}^{p \times p}$, where

$$H_{c(w)}[i, j] = \frac{\partial^2 c}{\partial w_i \partial w_j}(w)$$

E.g. $c(w) = \frac{1}{2}(w_1 + xw_2 - y)^2$, $\nabla c(w) = \begin{bmatrix} (w_1 + xw_2 - y) \\ (w_1 + xw_2 - y)x \end{bmatrix} = \begin{bmatrix} \frac{\partial c}{\partial w_1}(w) \\ \frac{\partial c}{\partial w_2}(w) \end{bmatrix}$.

$$H_{c(w)} = \begin{bmatrix} \frac{\partial}{\partial w_1} \frac{\partial}{\partial w_1} c(w) & \frac{\partial}{\partial w_1} \frac{\partial}{\partial w_2} c(w) \\ \frac{\partial}{\partial w_2} \frac{\partial}{\partial w_1} c(w) & \frac{\partial}{\partial w_2} \frac{\partial}{\partial w_2} c(w) \end{bmatrix} = \begin{bmatrix} 1 & x \\ x & x^2 \end{bmatrix}$$

For $(x = 0.1, y = 2)$ and $w_0 = (0, 0)$,

$$\nabla c(w) = \begin{bmatrix} -2 \\ -0.2 \end{bmatrix}, H_{c(w_0)} = \begin{bmatrix} 1 & 0.1 \\ 0.1 & 0.01 \end{bmatrix}$$

The eigenvalues of $H_{c(w_0)}$ are $\lambda_1 \approx 1, \lambda_2 = 0$. So we get a concave up in one direction, and flat in the other direction (a taco shape).

E.g. $c(w) = \frac{1}{2}(w_1 + x_1 w_2 - y_1)^2 + \frac{1}{2}(w_1 + x_2 w_2 - y_2)^2$. At $(x_1 = 0.1, y_1 = 2)$ and $(x_2 = -0.3, y_2 = -1)$,

$$H_{c(w_0)} = \begin{bmatrix} 2 & -0.2 \\ -0.2 & 0.1 \end{bmatrix}$$

for $w_0 = (0, 0)$. The eigenvalues are $\lambda_1 = 6.5, \lambda_2 = 0.08$.

Better Steps: $H_{c(w)}^{-1}$ is a matrix stepsize. The second-order method is

$$w_{t+1} = w_t - H_{c(w_t)}^{-1} \nabla c(w_t)$$

The scalar version was $w_{t+1} = w_t - \frac{1}{c''(w_t)} \cdot c'(w_t)$.

Using eigenvalue decomposition, $H_{c(w_t)}^{-1} \nabla c(w_t) = U \Lambda^{-1} \underbrace{U^T \nabla c(w_t)}_{\tilde{g}_t}$. This shows that the step is inversely proportional to curvature magnitude (λ_i 's). One step size could be bad, since in regular

gradient descent uses stepsize η , so Hessian allows us not to step too much on step surfaces and more on flat surfaces.

Computational issues: Computing $H_{c(w)}^{-1}$ is $O(p^2n)$ for linear regression. This is because $H_{c(w)} = \frac{1}{n} \sum_{i=1}^n H_{c_i(w)}$, where $H_{c_i(w)}$ takes $O(p^2)$ and you sum n times. Computing the inverse is $O(p^3)$, but for $p < n$, $p^3 < p^2n$, so $O(p^2n)$ overall.

An alternative is to use a vector of stepsizes $\eta_t \in \mathbb{R}^p$. So

$$H_{c(w_t)}^{-1} \approx \text{diag}(\eta_{t,1}, \eta_{t,2}, \dots, \eta_{t,p})$$

and the update is

$$w_{t+1} = w_t - \eta_t \underbrace{\cdot}_{\text{element-wise product}} \nabla c(w_t)$$

We do not use a fixed $\eta \in \mathbb{R}^p$. We use vector stepsizes, that change with time.

Definition 2.1: Adagrad

$$\eta_{t,j} = (1 + \bar{g}_{t,j})^{-1/2}$$

where $\bar{g}_{t,j} = \bar{g}_{t-1,j} + g_{t,j}^2$ and $\bar{g}_0 = 0$.

This has nice theory, but in practice, stepsizes shrink too much since we have a monotonic decrease in stepsize.

An intuitive modification is to take an exponential average of $g_{t,j}^2$, which leads to RMSProp.

Definition 2.2: RMSProp

$$v_{t,j} = (1 - \beta)g_{t,j}^2 + \beta v_{t-1,j}, \beta > 0$$

$$\eta_{t,j} = \frac{\eta}{\sqrt{v_{t,j} + \varepsilon}} = \eta(v_{t,j} + \varepsilon)^{-1/2}$$

Also add **momentum**: with fixed stepsize η ,

$$w_{t+1} = w_t - \eta m_{t+1}, m_{t+1} = g_t + \beta_2 m_t, \beta_2 > 0$$

An alternate form is

$$w_{t+1} = w_t - \eta g_t + \beta(w_t - w_{t-1})$$

Definition 2.3: Adam

Puts RMSProp and momentum together with extra modification.

There is a bias correct for $v_{0,j} = 0$. There is an exponential average in momentum $m_{t+1} = (1 - \beta_2)g_t + \beta_2 m_t$. The last modification is

$$\text{Adam's } \frac{1}{\sqrt{v_{t,j}} + \varepsilon} \text{ vs. RMSProp's } \frac{1}{\sqrt{v_{t,j} + \varepsilon}}$$

2.2 Comparisons

Second-order updates takes $O(p^2n)$ per iteration and gradient descent takes $O(pn)$ per iteration.

Definition 2.4: Stochastic Gradient Descent (SGD)

Randomly pick $\mathcal{B}_t \subseteq \{1, \dots, n\}$ with $|\mathcal{B}_t| = b$, the gradient $\hat{g}_t = \frac{1}{b} \sum_{i \in \mathcal{B}_t} \nabla c_i(w_t)$,

$$w_{t+1} = w_t - \eta_t \hat{g}_t$$

The runtime is $O(pb)$.

This is much faster than $O(pn)$ for GD. $\hat{g}_t \approx g_t = \frac{1}{n} \sum_{i=1}^n c_i(w_t)$.